

Training an ML Model



The Dataset

So far, we have introduced our simulator

The rest of our plan is as follows

- We learn an ML model
- We embed the model in a larger optimization problem
- We obtain a solution, i.e. a set of action to control the epidemics

But which data are we going to use for training?



The Dataset

So far, we have introduced our simulator

The rest of our plan is as follows

- We learn an ML model
- We embed the model in a larger optimization problem
- We obtain a solution, i.e. a set of action to control the epidemics

But which data are we going to use for training?

Since we have a simulator, we can **build our dataset**

- This means we can generate as much data as we wish
- ...But also that we are responsible for **how to generate it**



Building Our Dataset

We need to define the **structure** of the dataset

- We will focus on **Non-Therapeutic Interventions** (mask mandates, etc.)
- NPIs affect the **β parameter** in a SIR model (γ will be fixed)
- We will focus on making predictions at **weekly intervals**

Therefore, we can cover our needs with...

- For the **input**:
 - The state (S_0, I_0, R_0) at the beginning of a week
 - The value of β , i.e. a proxy for the NPIs currently in effect
- For **output**: the state after one week (S, I, R)

Given an input (S, I, R, β) , we can get the output via simulation



Which input configurations should we generate?



Building Our Dataset

A training set should be **representative of the test distribution**

- We do not have a fixed test distribution (no test set)
- ...But we know that the ML model will be **used by an optimizer**

The optimizer will seek to minimize the total infections. So, we will need:

- High accuracy on the best configurations, so as to **find** them
- High accuracy on the worst configurations, so as to **avoid** them

I.e. to be safe the model should work **all across the board**

Hence, we need to cover reasonably well all the input space

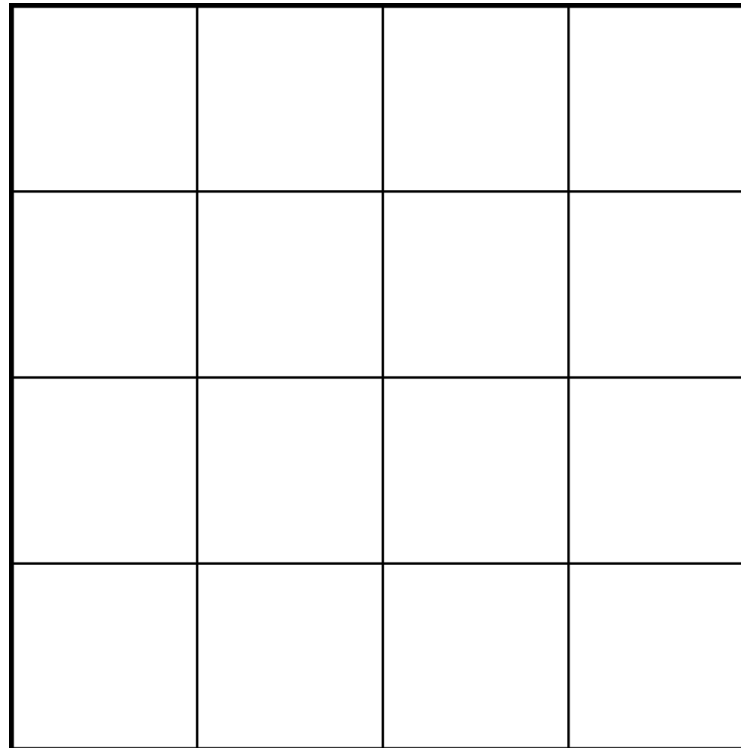
- The simplest approach would be use use a regular grid (**factorial design**)
- ...But factorial design has poor scalability



Latin Hypercube Sampling

The method we will use is called **Latin Hypercube Sampling**

Suppose we want to sample m points for n attributes with fixed ranges



- We can view the sampling space as a hypercube
- ...Then we divide each dimension in n segments

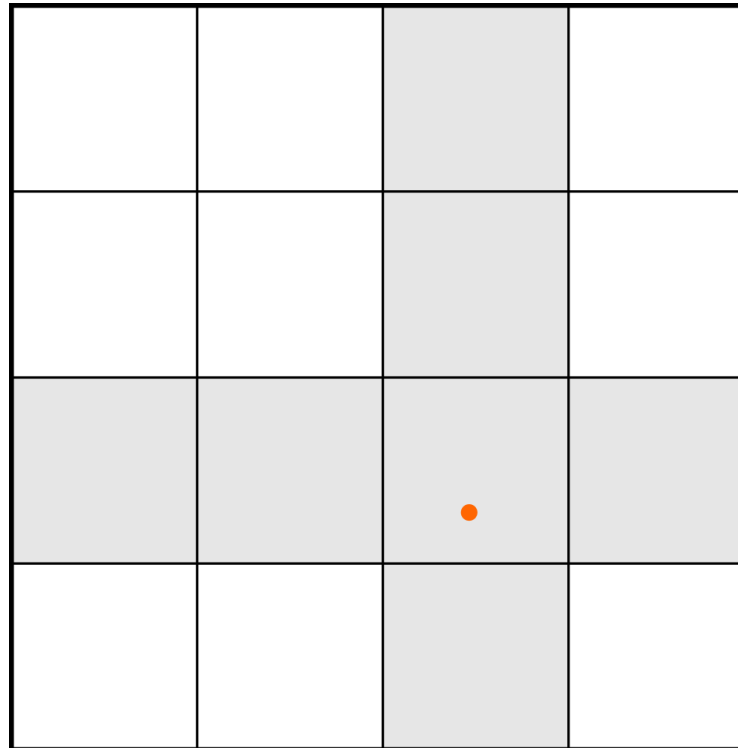
In the example we will to sample 4 points for 2 attributes



Latin Hypercube Sampling

The method we will use is called **Latin Hypercube Sampling**

Suppose we want to sample m points for n attributes with fixed ranges



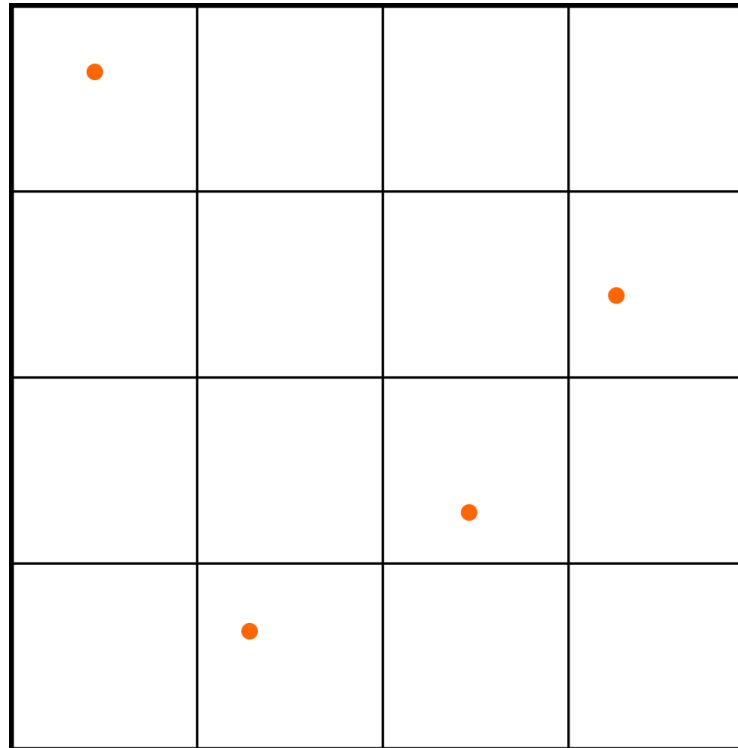
- We sample the first point uniformly at random
- ...Then we **cover** the row and column that contain the sample



Latin Hypercube Sampling

The method we will use is called **Latin Hypercube Sampling**

Suppose we want to sample m points for n attributes with fixed ranges



- When we take additional samples, we **exclude all covered row/columns**
- ...So we end up with a pattern similar to that of the figure

LHS can cover quite uniformly a given space with relatively few samples

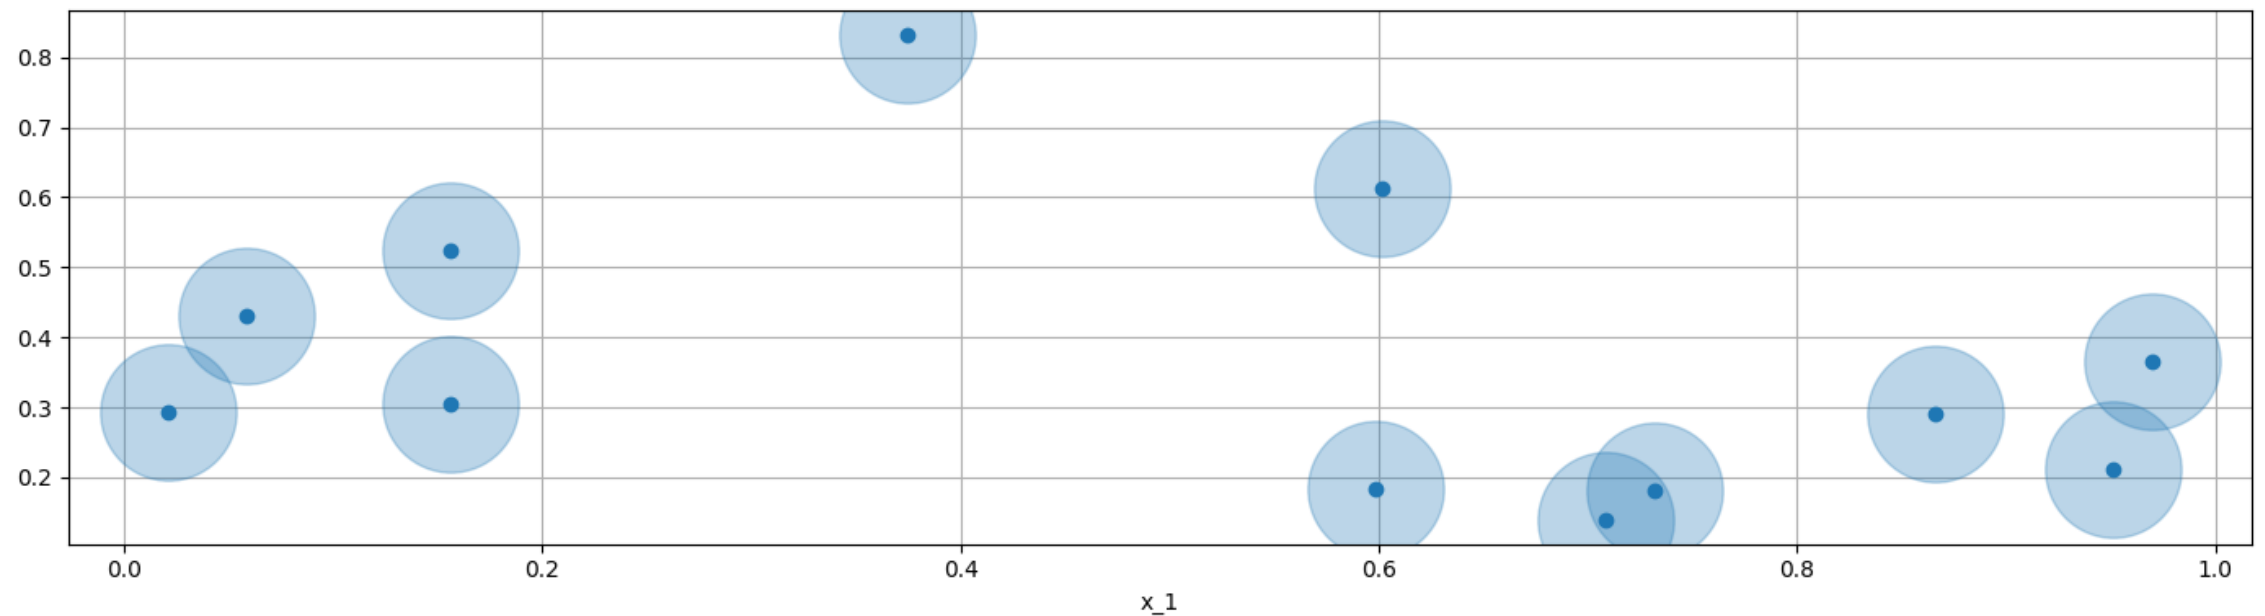


Latin Hypercube Sampling

Let's see a practical example

Here is the result of **uniform sampling**, for reference

```
In [2]: test_nsamples, test_ranges = 12, [(0., 1.), (0., 1.)]  
X = util.sample_points(test_ranges, test_nsamples, mode='uniform', seed=42)  
util.plot_2D_samplespace(X, figsize=figsize)
```

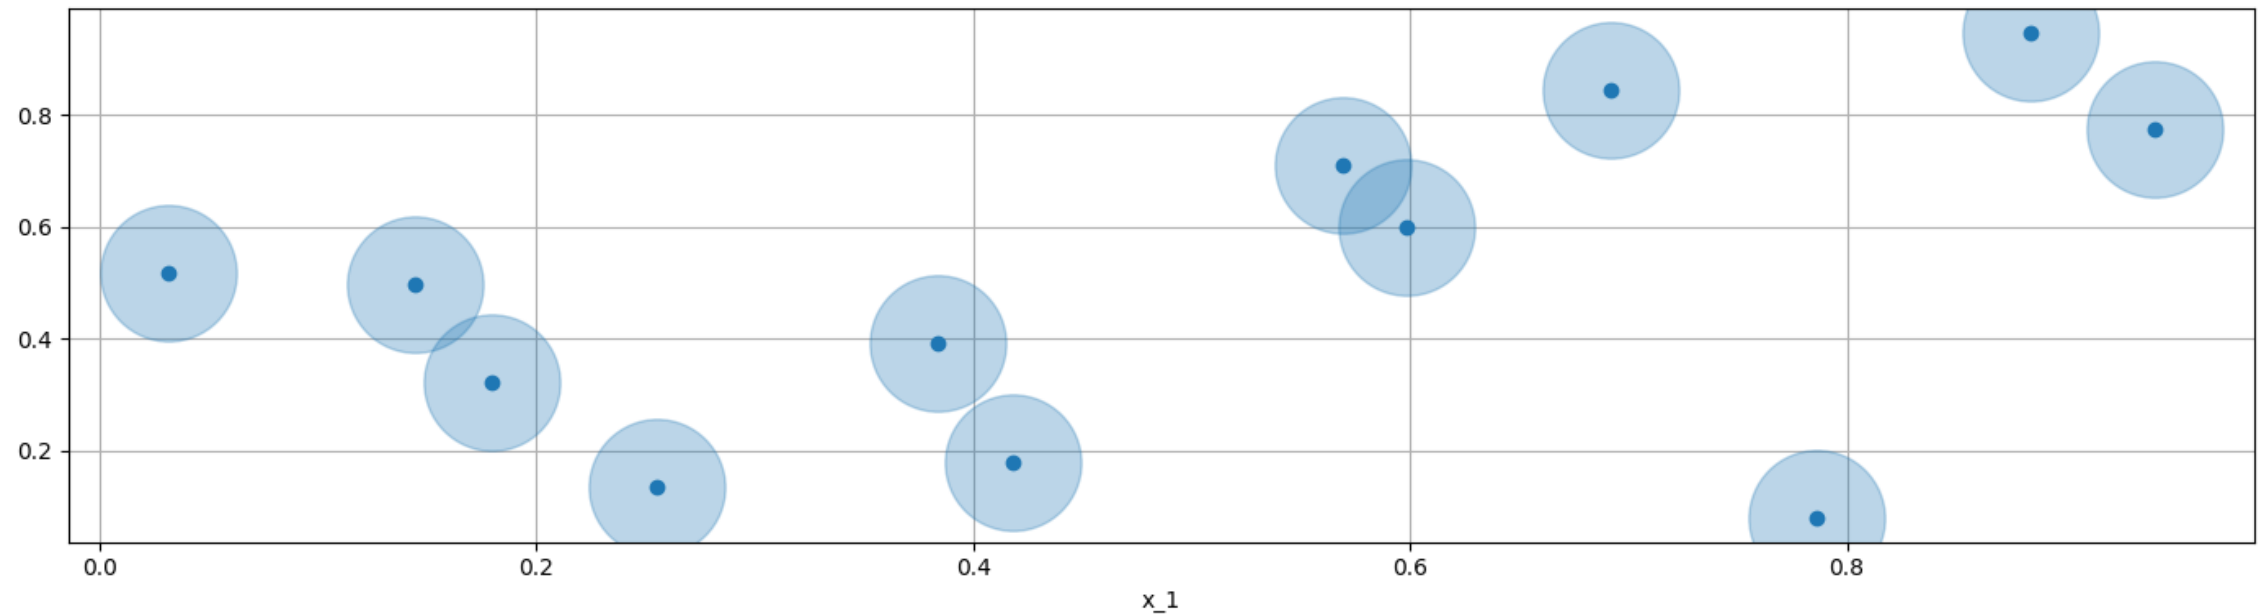


Latin Hypercube Sampling

Let's see a practical example

...And here is the result of classical LHS:

```
In [3]: test_nsamples, test_ranges = 12, [(0., 1.), (0., 1.)]  
X = util.sample_points(test_ranges, test_nsamples, mode='lhs', seed=42)  
util.plot_2D_samplespace(X, figsize=figsize)
```

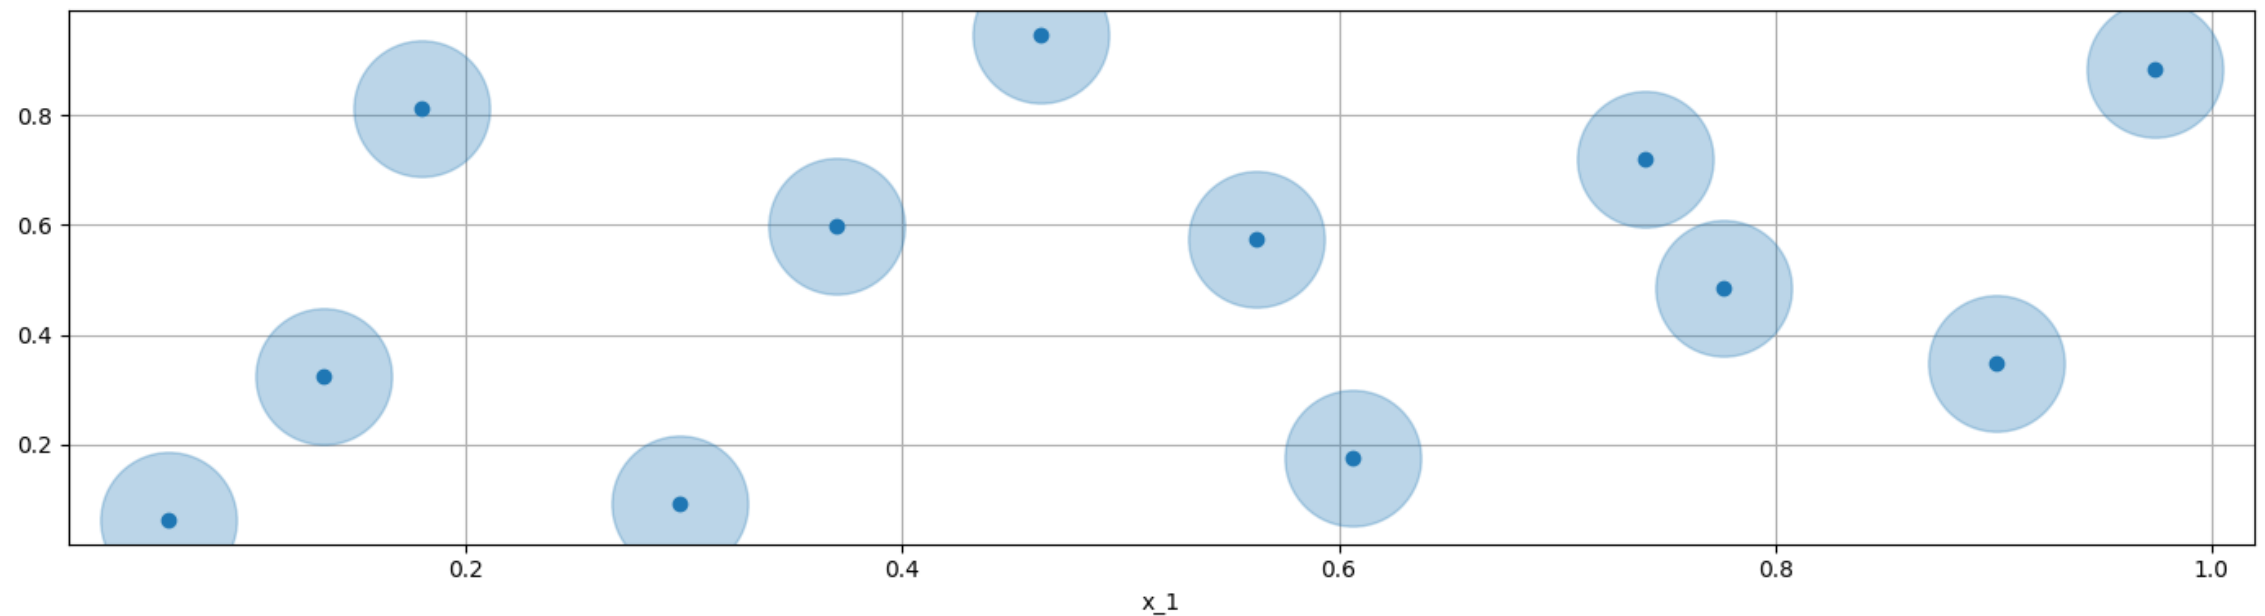


Latin Hypercube Sampling

The process can be further improved

E.g. we can sample multiple times to **maximize the minimum** distance

```
In [4]: test_nsamples, test_ranges = 12, [(0., 1.), (0., 1.)]  
X = util.sample_points(test_ranges, test_nsamples, mode='max_min', seed=42)  
util.plot_2D_samplespace(X, figsize=figsize)
```



Dataset Input

We are now ready to generate our **dataset input**

```
In [5]: n_tr, n_ts = 10000, 2000
        sir_tr_in = util.generate_SIR_input(max_samples=n_tr, mode='max_min', seed=42, normalize=True, r
        sir_ts_in = util.generate_SIR_input(max_samples=n_ts, mode='max_min', seed=42, normalize=True, r
        sir_tr_in.head()
```

Out [5]:

	S	I	R	beta
0	0.352572	0.512310	0.135119	0.167925
1	0.175114	0.707867	0.117019	0.100628
2	0.024010	0.836985	0.139005	0.048884
3	0.139277	0.247573	0.613150	0.367367
4	0.491191	0.195846	0.312963	0.067067

- We sample S, I, R, β from $[0, 1]^3 \times [0, .4]$
- ...Then S, I, R are normalized so that their sum is 1

This will reduce in some redundancy in the dataset



Dataset Output

We obtain the corresponding output via simulation

```
In [6]: %%time
gamma = 1/14
sir_tr_out = util.generate_SIR_output(sir_tr_in, gamma, 7)
sir_ts_out = util.generate_SIR_output(sir_ts_in, gamma, 7)
sir_tr_out.head()
```

```
CPU times: user 4.2 s, sys: 6.64 ms, total: 4.21 s
Wall time: 4.21 s
```

Out[6]:

	S	I	R
0	0.201814	0.425756	0.372430
1	0.115945	0.474359	0.409696
2	0.019150	0.511369	0.469481
3	0.078295	0.196566	0.725139
4	0.453265	0.148189	0.398546

- We picked $\gamma = 1/14$ (this will be fixed in our use case)
- For each configuration, we simulate one week

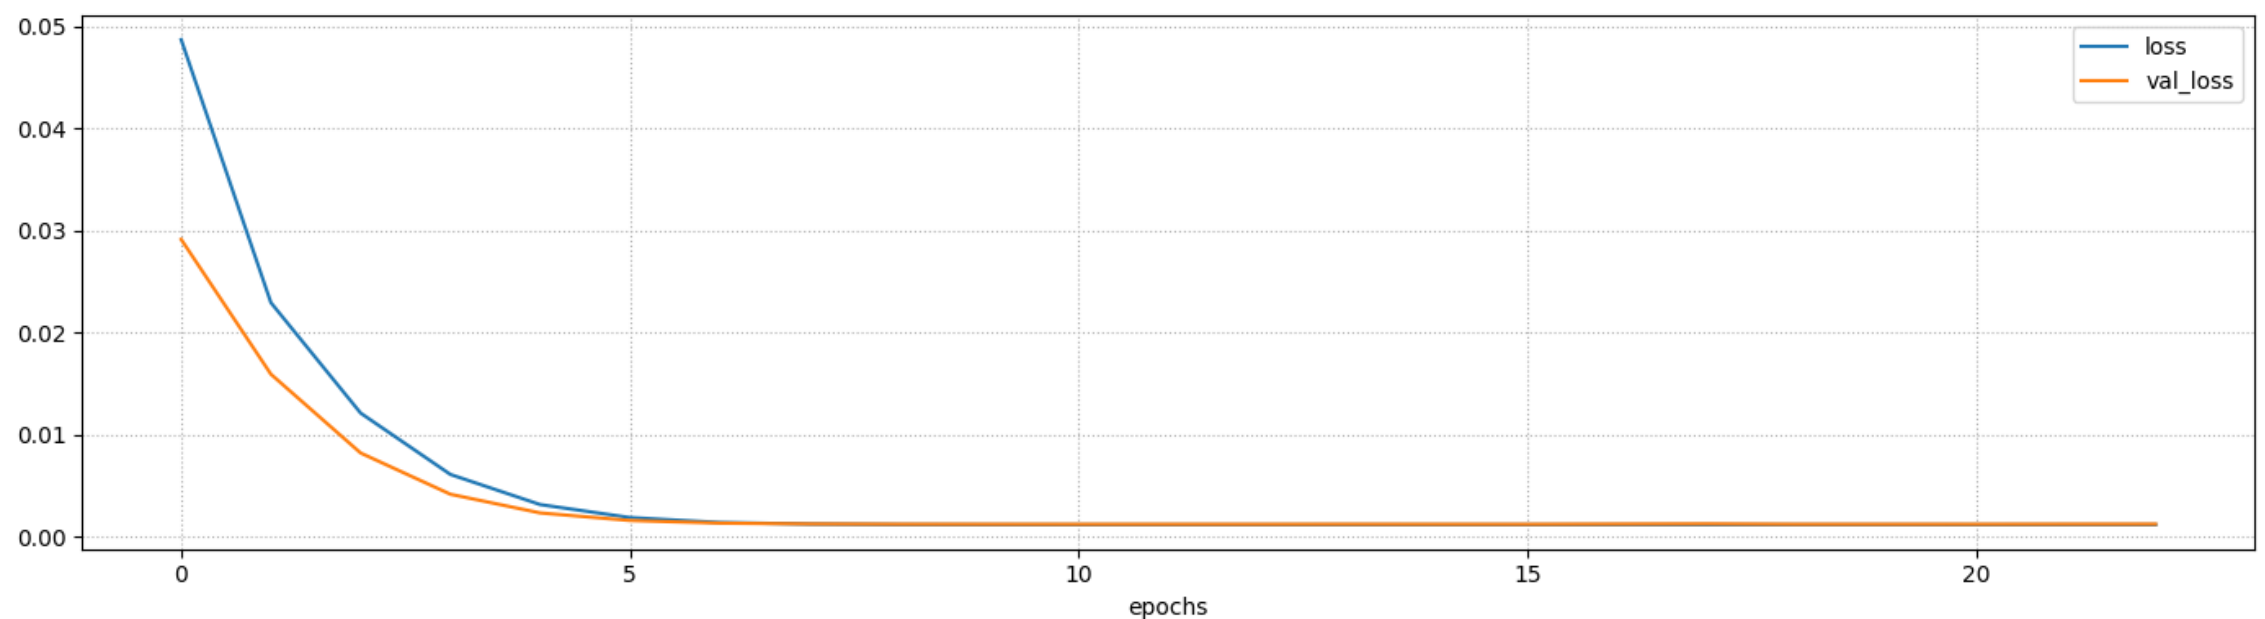


Training a Model

We try with Linear Regression

```
In [8]: nn0 = util.build_nn_model(input_shape=4, output_shape=3, hidden=[], name='LR')
history0 = util.train_nn_model(nn0, sir_tr_in, sir_tr_out, loss='mse', verbose=0, validation_split=0.1)
util.plot_training_history(history0, figsize=figsize)
util.print_ml_metrics(nn0, sir_tr_in, sir_tr_out, 'training')
util.print_ml_metrics(nn0, sir_ts_in, sir_ts_out, 'test')
```

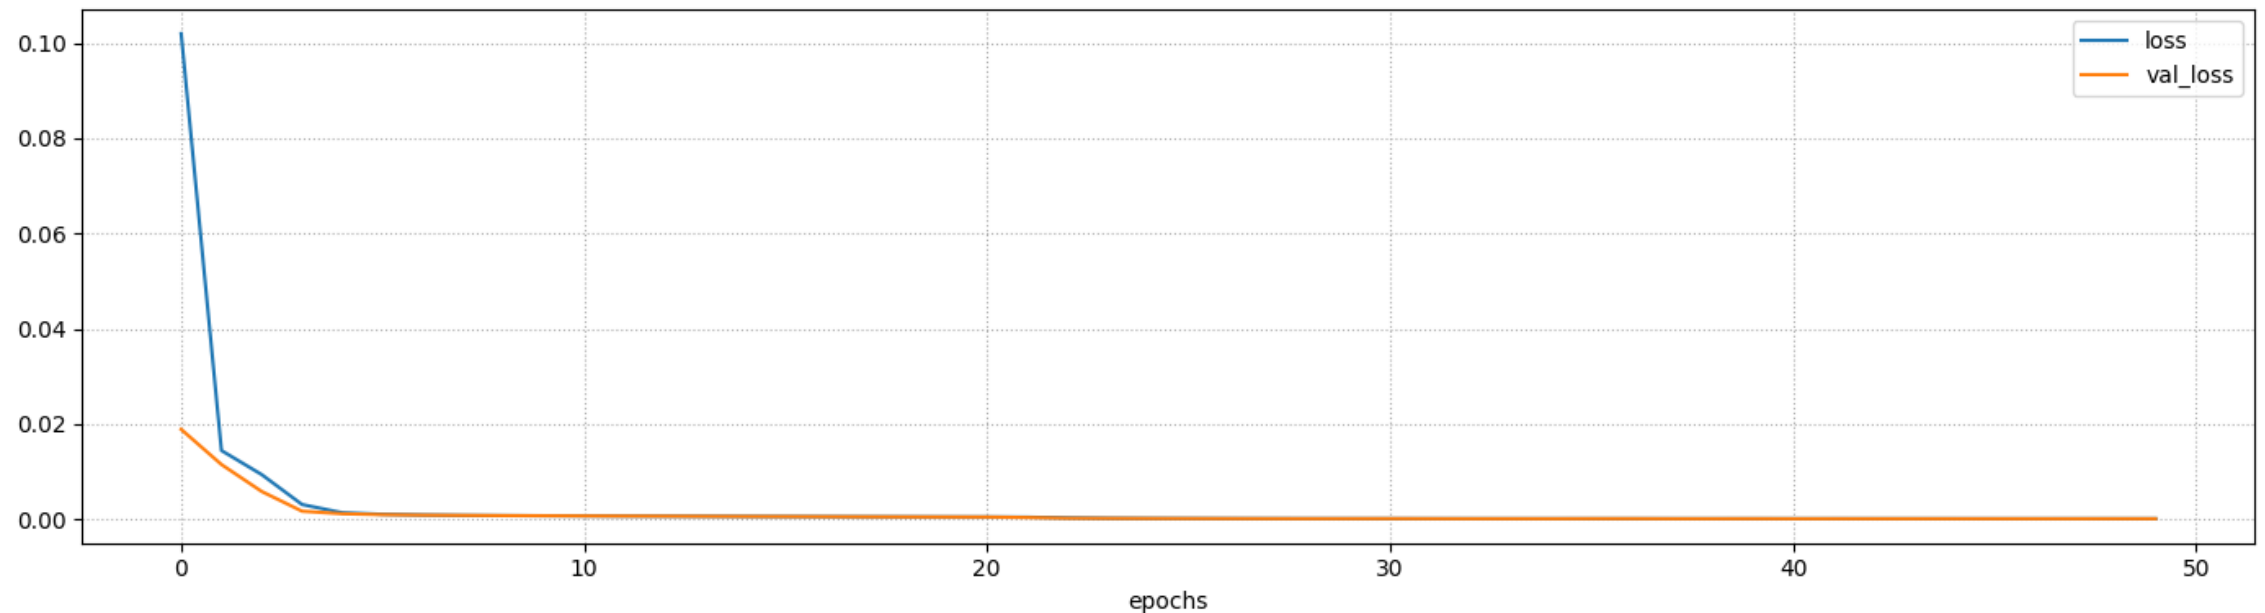
2022-12-12 17:06:20.411632: I tensorflow/core/platform/cpu_feature_guard.cc:193] This TensorFlow binary is optimized with oneAPI Deep Neural Network Library (oneDNN) to use the following CPU instructions in performance-critical operations: AVX2 FMA
To enable them in other operations, rebuild TensorFlow with the appropriate compiler flags.



Training a Model

...And with a shallow Neural Network

```
In [9]: nn1 = util.build_nn_model(input_shape=4, output_shape=3, hidden=[8], name='MLP')
history1 = util.train_nn_model(nn1, sir_tr_in, sir_tr_out, loss='mse', verbose=0, validation_split=0.1)
util.plot_training_history(history1, figsize=figsize)
util.print_ml_metrics(nn1, sir_tr_in, sir_tr_out, 'training')
util.print_ml_metrics(nn1, sir_ts_in, sir_ts_out, 'test')
```



Final loss: 0.0001 (training), 0.0001 (validation)

R2: 0.99, MAE: 0.0072, RMSE: 0.01 (training)

R2: 0.99, MAE: 0.0072, RMSE: 0.01 (test)



Considerations and Next Steps

We will save both models for later

```
In [10]: util.save_ml_model(nn0, 'nn0')  
         util.save_ml_model(nn1, 'nn1')
```

- The network is much better in terms of accuracy
- ...But the Linear Regressor is simpler!

Hence, the approaches provide different trade offs

We are halfway there

We now have our ML model(s)!

- We need to understand how they can be embedded in an optimization model
- ...And we need to define our optimization model itself

