

Encoding ML Models



Encoding ML Models

We need to embed our ML model into an optimization model

The basic approach is based on two observations:

- A neural network is a collection of connected neurons
- So we just need to **encode each neuron** using a given optimization method

Let's consider for example a ReLU neuron

$$y = \max(0, \boldsymbol{w}^T \boldsymbol{x} + \theta)$$

Where \boldsymbol{w} is the (row) vector of weights and \boldsymbol{b} is the bias. It can be encoded by:

- Introducing a variable for each input
- Introducing a variable for the output
- Modeling (e.g. in MILP, or SMT, or CP) the sum and max operators



Encoding ML Models

In this case, we will adopt a MILP encoding for the relation
 $y = \max(0, wx + \theta)$

$$y - s = wx + \theta$$

$$z = 1 \Rightarrow s \leq 0$$

$$z = 0 \Rightarrow y \leq 0$$

$$y, s \geq 0, x \in \mathbb{R}^n, z \in \{0, 1\}$$

- s is an auxiliary **slack variable** and z is an auxiliary binary variable
- The implications are called **indicator constraints**
- They are handled natively by some MILP solvers, or they can be linearized



Encoding ML Models

Let's have a better look at the encoding:

$$y - s = wx + \theta$$

$$z = 1 \Rightarrow s \leq 0$$

$$z = 0 \Rightarrow y \leq 0$$

$$y, s \geq 0, x \in \mathbb{R}^n, z \in \{0, 1\}$$

If $z = 1$, it means that the neuron is **active**

- In this case s is forced to 0, we have: $y = wx + \theta$
- ...And $wx + \theta$ is non-negative



Encoding ML Models

Let's have a better look at the encoding:

$$y - s = wx + \theta$$

$$z = 1 \Rightarrow s \leq 0$$

$$z = 0 \Rightarrow y \leq 0$$

$$y, s \geq 0, x \in \mathbb{R}^n, z \in \{0, 1\}$$

If $z = 0$, it means that the neuron is **inactive**

- In this case y (the neuron output) is 0
- We have $s = wx + \theta$ (note that s does not contribute to the neuron output)
- ...And $wx + \theta$ is negative



Loading the Network

We will handle the encoding via the EMLlib

It's a small (and still rough) library for the EML approach. It allows to:

- Load models from ML libraries (currently NNs from keras, DTs from sklearn)
- Convert them into an internal format
- ...And translate such format into an encoding for a target technique

As a first step, let's load both our trained models

```
In [35]: knn0 = util.load_ml_model('nn0')  
         knn1 = util.load_ml_model('nn1')
```

...And then convert the in the EMLlib internal format:

```
In [36]: nn0 = keras_reader.read_keras_sequential(knn0)  
         nn1 = keras_reader.read_keras_sequential(knn1)
```



Loading the Network

Printing the loaded network shows "bounds" for each neuron

This is easier to parse for the Linear Regression model:

```
In [37]: nn0
```

```
Out[37]: [input] (0, 0):[-inf, inf] (0, 1):[-inf, inf] (0, 2):[-inf, inf] (0, 3):[-inf, inf]
          [dense,linear] (1, 0):[-inf, inf]/[-inf, inf] (1, 1):[-inf, inf]/[-inf, inf] (1, 2):[-inf, in
          f]/[-inf, inf]
```

The bounds represent the **domain** of output of each neuron

- Currently, most output have an **infinite range**

This is a problem since our MILP encoding for each neuron...

...Requires finite bounds to linearize the indicator constraints

- Actually, the tighter the bounds, the better the MILP encoding will work



Loading the Network

We can use 0 and 1 as bounds for all our inputs

...Since the population is normalized and β is typically a low value

```
In [38]: nn0.layer(0).update_lb(np.zeros(4))  
nn0.layer(0).update_ub(np.ones(4));
```

Internal bounds can be inferred with one round of **constraint propagation**

- This needs to be done for the weighted sum in each neuron
- ...And for the ReLU, when actually present

The process is implemented in the `ibr_bounds` function:

```
In [43]: from eml.net.process import ibr_bounds  
ibr_bounds(nn0)  
nn0
```



Network and I/O Variables

We will consider a planning problem over eah weeks

So, we will create variables to represent S , I , R , and β for those weeks:

$$\begin{aligned}\beta_t &\in [0, 1] & \forall t = 0..eah - 1 \\ S_t, I_t, R_t &\in [0, 1] & \forall t = 0..eah\end{aligned}$$

- For week t we do not need to represent the value of β
- ...But we need the S , I , R variables to represent the final state
- S_0, I_0, R_0 will be fixed to the values from the initial state

Our objective will be to maximize S_{eah}



Network and I/O Variables

We will encode **an instance** of our network **for each week**

Each will connect consecutive values of S, I, R

$$\begin{aligned}(S_{t+1}, I_{t+1}, R_{t+1}) &= \text{NN}(S_t, I_t, R_t, \beta_t) & \forall t = 0..eoh - 1 \\ \beta_t &\in [0, 1] & \forall t = 0..eoh - 1 \\ S_t, I_t, R_t &\in [0, 1] & \forall t = 0..eoh\end{aligned}$$

- $\text{NN}(\dots)$ represents the network encoding
- ...i.e. the equations we discussed early on in the lecture
- Each network takes as input the value of S, I, R , and β for week t
- ...And links them to the value of S, I, R for week $t + 1$



Network ad I/O Variables

The code for the planning problem is in `solve_sir_planning`

We use the CBC solver via Google Or-tools:

```
slv = pywraplp.Solver.CreateSolver('CBC')
```

We start by building the network I/O variables:

```
for t in range(nweeks+1):  
    X['S', t] = slv.NumVar(0, 1, f'S_{t}')  
    X['I', t] = slv.NumVar(0, 1, f'I_{t}')  
    X['R', t] = slv.NumVar(0, 1, f'R_{t}')  
    if t < nweeks: X['b', t] = slv.NumVar(0, 1, f'b_{t}')
```

- The network will be embedded as an **encoding**
- ...Which cannot be defined unless we have the variables first



Network Encodings

The library we use handles multiple solvers via "backend" objects

Therefore we need to build a backend for Or-tools:

```
bkd = ortools_backend.OrtoolsBackend()
```

- The backend defines the primitives to build the NN constraints

The encoding themselves are built using the `encode` function:

```
for t in range(1, nweeks+1):  
    vin = [X['S',t-1], X['I',t-1], X['R',t-1], X['b',t-1]]  
    vout = [X['S',t], X['I',t], X['R',t]]  
    encode(bkd, nn, slv, vin, vout, f'nn_{t}')
```

- Neurons are processed one by one
- Intermediate variables are built as needed



Non-Pharmaceutical Intervention

Now we need to setup the rest of the optimization model

...Since we delayed this even too much to focus on the NN encoding

- At each week we can choose to apply a number of **NPIs**
- ...Which (we remind) stands for "Non-Pharmaceutical Interventions"

We will assume each NPI i has a (socio-economical) **cost c_i**

- ...And can reduce the current β value by a **factor r_i**
- β has a "base value", which depends on the disease itself

So, if we apply NPIs 1, 3, and 4:

- We pay a cost equal to $c_1 + c_3 + c_4$
- And we have $\beta = r_1 r_3 r_4 \beta_{base}$

Using multiple NPIs has diminishing returns



Non-Pharmaceutical Intervention

This part of the problem can be formalized as follows:

We introduce a binary variable x_{it} for each NPI and week (except the last)

$$x_{it} \in \{0, 1\} \quad \forall i = 1 \dots n_{npi}, \forall t = 0 \dots eoh - 1$$

- $x_{it} = 1$ iff we apply NPI i at week t

We assume the total cost should not exceed a given budget

$$\sum_{t=0}^{eoh-1} \sum_{i=1}^{n_{npi}} c_i x_{it} \leq C$$

- Where C is the budget value



Non-Pharmaceutical Intervention

The effect on β is non-linear and trickier to handle

We linearize it by introducing multiple variables for β at each week

- β_{0t} represents the "base" β value
- β_{it} represents β as affected by the i -th NPI
- Therefore $\beta_{n_{npi},t}$ is the same as the variable connected to the NN for week t

For each intermediate variable we have:

$$\beta_{it} \geq r_i \beta_{i-1,t} - 1 + x_{it} \quad \forall i = 1..n_{npi}, \forall t = 0..eoh - 1$$

$$\beta_{it} \geq \beta_{i-1,t} - x_{it} \quad \forall i = 1..n_{npi}, \forall t = 0..eoh - 1$$

- If $x_{it} = 1$, the first constraint is active and the second is trivialized
- If $x_{it} = 0$, the opposite is true



Non-Pharmaceutical Intervention

An analogous set of constraints handles the upper bounds

$$\beta_{it} \leq r_i \beta_{i-1,t} + 1 - x_{it} \quad \forall i = 1..n_{npi}, \forall t = 0..eoh - 1$$

$$\beta_{it} \leq \beta_{i-1,t} + x_{it} \quad \forall i = 1..n_{npi}, \forall t = 0..eoh - 1$$

Together with the previous set:

$$\beta_{it} \geq r_i \beta_{i-1,t} - 1 + x_{it} \quad \forall i = 1..n_{npi}, \forall t = 0..eoh - 1$$

$$\beta_{it} \geq \beta_{i-1,t} - x_{it} \quad \forall i = 1..n_{npi}, \forall t = 0..eoh - 1$$

...We obtain the desired behavior, e.g.:

- If $x_{1t} = x_{3t} = x_{4t} = 1$
- ...Then $\beta_{n_{npi}t} = r_1 r_3 r_4 \beta_{0,t}$

The details of the code can be found in the `solve_sir_planning`



Solving the NPI Planning Problem

We will consider the following set of NPIs

```
In [46]: npis = [  
    util.NPI('masks-indoor', effect=0.75, cost=1),  
    util.NPI('masks-outdoor', effect=0.9, cost=1),  
    util.NPI('dad', effect=0.7, cost=3),  
    util.NPI('bar-rest', effect=0.6, cost=3),  
    util.NPI('transport', effect=0.6, cost=4)  
]
```

We will consider a horizon of 3 weeks and the following parameters:

```
In [48]: S0, I0, R0 = 0.99, 0.01, 0.00  
nweeks = 3  
tlim = 30  
beta_base = 0.35  
budget = 20  
gamma = 1/14
```



Solution and Evaluation

Let's start by using the (much more accurate) NN model

```
In [55]: %%time
sol, closed = util.solve_sir_planning(knn1, npis, S0, I0, R0, beta_base=beta_base, budget=budget,
                                     nweeks=nweeks, tlim=tlim)

print(f'Problem closed: {closed}')
sol_df = util.sol_to_dataframe(sol, npis, nweeks)
sol_df
```

```
Problem closed: True
CPU times: user 1.57 s, sys: 60.8 ms, total: 1.63 s
Wall time: 1.63 s
```

Out[55]:

	S	I	R	b	masks-indoor	masks-outdoor	dad	bar-rest	transport
0	0.990000	0.010000	0.000000	0.14175	1.0	1.0	0.0	1.0	0.0
1	0.944169	0.014658	0.010808	0.09450	1.0	0.0	0.0	1.0	1.0
2	0.929573	0.010191	0.023377	0.11025	1.0	0.0	1.0	1.0	0.0
3	0.912024	0.017791	0.036351	NaN	NaN	NaN	NaN	NaN	NaN

■ The result seem reasonable

■ ...But how can we know for sure?



Solution and Evaluation

Our optimization model relies on **predictions**

We need to test their quality on the simulator:

```
In [57]: beta_sched = sol_df.iloc[:,-1]['b']  
util.simulate_SIR_NPI(S0, I0, R0, beta_sched, gamma, steps_per_day=1)
```

Out[57]:

	S	I	R
0	0.990000	0.010000	0.000000
1	0.977482	0.016106	0.006412
2	0.966344	0.018582	0.015075
3	0.950792	0.023622	0.025586

Unless we've been unlucky during training (it's stochastic!)

- The final value for S should be close to 0.95
- ...And **possibly quite different** from our model predictions!



Solution and Evaluation

It's even more clear if we use the Linear Regression model

```
In [58]: %%time
sol2, closed2 = util.solve_sir_planning(knn0, npis, S0, I0, R0, beta_base=beta_base, budget=budget,
                                         nweeks=nweeks, tlim=tlim)

print(f'Problem closed: {closed2}')
sol_df2 = util.sol_to_dataframe(sol2, npis, nweeks)
sol_df2
```

```
Problem closed: True
CPU times: user 487 ms, sys: 33.3 ms, total: 521 ms
Wall time: 521 ms
```

Out[58]:

	S	I	R	b	masks-indoor	masks-outdoor	dad	bar-rest	transport
0	0.990000	0.010000	0.000000	0.1575	1.0	0.0	0.0	1.0	0.0
1	0.765027	0.183567	0.053728	0.0945	1.0	0.0	0.0	1.0	1.0
2	0.608145	0.235826	0.160640	0.0945	1.0	0.0	0.0	1.0	1.0
3	0.492628	0.235837	0.278775	NaN	NaN	NaN	NaN	NaN	NaN

■ Now the solution process is **very fast**

■ ...And it **looks like a disaster**



Solution and Evaluation

However, if we evaluate the solutions via the simulator...

```
In [59]: beta_sched2 = sol_df2.iloc[:,-1]['b']  
util.simulate_SIR_NPI(S0, I0, R0, beta_sched2, gamma, steps_per_day=1)
```

Out[59]:

	S	I	R
0	0.990000	0.010000	0.000000
1	0.975265	0.017934	0.006801
2	0.962909	0.020653	0.016438
3	0.948935	0.023577	0.027488

...They are not bad at all!

Our ML models are making mistakes

- For many reasons: bias, compound error, "weak spots"
- ...But as long as they **guide the solver** in the right place, we get a good solution

This is good news, but leaves some open issues (see later)



Solution and Evaluation

The main issue is: how much can we trust our models?

In our case, it turns out the answer is "a lot"

■ Here's what we get by solving the problem via brute force:

```
In [62]: %%time
best_S, best_sched = util.solve_sir_brute_force(npis, S0, I0, R0, beta_base, gamma, nweeks, budg
best_S

CPU times: user 25.9 s, sys: 0 ns, total: 25.9 s
Wall time: 25.9 s

Out[62]: 0.9554715100410379
```

- The NN solution in particular is actually pretty good
- ...And we obtain it in much less time
- As the problem size grows, the gap in computation time becomes larger



Some Considerations

This kind of hybrid approach can be complex to build

- But sometimes it's (almost) the only choice!
- It generally worked in our case

EML-like approaches can be used to generate adversarial examples

- It is at the basis of some tools for NN verification

There are several open issues

- The optimizer often ends up finding weaknesses in the ML model
- The approach scalability is limited

